

Movie Data Analysis

This notebook explores various aspects of movie data, including revenue, profit margins, release windows, and more. It utilizes visualizations and statistical techniques to uncover trends and patterns in the data.

Business Understanding

In this project, I analyzed historical movie data to uncover actionable insights for a new movie studio venture. My goal was to identify the critical factors that drive the commercial success of movies, focusing on profitability, audience engagement, and production efficiency. By leveraging data science techniques, I aimed to enable data-driven decisions on budget allocation, genre prioritization, and release strategies, ultimately maximizing return on investment (ROI) while minimizing financial risks.

Key Business Questions:

- Which **genres** generate the highest revenue, and should the studio invest in them?
- How do **production budgets** correlate with worldwide gross, and can a higher budget ensure success?
- What are the most **profitable release windows** for maximizing revenue?
- Which **studios** achieve the best profit margins, and what can be learned from them?
- How does the **original language** of a movie influence its global performance?
- Can specific **movie attributes** be used to predict revenue and inform production strategies?

Dataframes for Analysis

To address the business questions and achieve the project goals, I identified and utilized the following datasets:

1. Box Office Mojo Data

- **Key Variables:** title , studio , domestic_gross , foreign_gross , year
- **Purpose:** I used this data to evaluate box office trends, compare domestic and international earnings, and assess studio performance over time.

2. The Numbers Data

- **Key Variables:** movie , production_budget , domestic_gross , worldwide_gross , release_date
- **Purpose:** This dataset helped me analyze the relationship between production

budgets and revenue to identify profitability patterns.

3. Rotten Tomatoes Movie Info Data

- **Key Variables:** rating , genre , director , runtime , box_office
- **Purpose:** I explored how genre, director influence, runtime, and audience ratings impacted a movie's box office performance.

4. TheMovieDB Data

- **Key Variables:** title , popularity , vote_average , vote_count , release_date
- **Purpose:** I examined the correlation between popularity metrics (e.g., vote counts, average ratings) and box office success to understand audience engagement.

5. IMDb Database (im.db.zip)

- **Key Tables:** movie_basics , movie_ratings
- **Purpose:** Using SQLite, I queried this database to integrate movie attributes (e.g., genres, languages, and ratings) with performance data for deeper insights.

By combining these datasets, I ensured a comprehensive analysis of movie performance from multiple perspectives, aligning with the project objectives.

Data preparation

Libraries and Data Import

```
In [1]: # Import libraries for data manipulation and visualization
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Import libraries for file handling and database interaction
import gzip
import sqlite3
import os

# Set default visualization style
sns.set_theme(style="whitegrid")
sns.set_palette("pastel")
plt.rcParams['figure.figsize'] = [10, 6]
```

2.1. Box Office Mojo Data

```
In [2]: # Define the path
file_path = 'C:/Users/USER/Desktop/Movie-Project/data/raw/zippedData/bom.movie_gros

try:
    # Load the gzipped CSV directly
```

```
bom_gross = pd.read_csv(file_path, compression='gzip')

# Display general information about the dataframe
print("Dataframe Info:")
bom_gross.info()

# Display the first few rows
print("\nFirst few rows of data:")
display(bom_gross.head())

# Show data types of each column
print("\nColumn Data Types:")
print(bom_gross.dtypes)

except Exception as e:
    print(f"Error loading file: {e}")
```

Dataframe Info:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3387 entries, 0 to 3386
Data columns (total 5 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   title                 3387 non-null   object
 1   studio                3382 non-null   object
 2   domestic_gross        3359 non-null   float64
 3   foreign_gross         2037 non-null   object
 4   year                  3387 non-null   int64
dtypes: float64(1), int64(1), object(3)
memory usage: 132.4+ KB
```

First few rows of data:

	title	studio	domestic_gross	foreign_gross	year
0	Toy Story 3	BV	415000000.0	652000000	2010
1	Alice in Wonderland (2010)	BV	334200000.0	691300000	2010
2	Harry Potter and the Deathly Hallows Part 1	WB	296000000.0	664300000	2010
3	Inception	WB	292600000.0	535700000	2010
4	Shrek Forever After	P/DW	238700000.0	513900000	2010

Column Data Types:

```
title          object
studio         object
domestic_gross float64
foreign_gross  object
year           int64
dtype: object
```

2.2 The Numbers Data

```
In [3]: # Load The Numbers (movie budgets) dataset
file_path = 'C:/Users/USER/Desktop/Movie-Project/data/raw/zippedData/tn.movie_budge
```

```
try:
    tn_budgets = pd.read_csv(file_path, compression='gzip')

    print("The Numbers Data Preview:")
    display(tn_budgets.head()) # Display the first few rows

    print("\nDataset Info:")
    tn_budgets.info() # Get an overview of the dataset

    # Check for duplicate rows
    print("\nNumber of duplicate rows:", tn_budgets.duplicated().sum())

    print("\nMissing values per column:")
    print(tn_budgets.isnull().sum())

    print("\nSummary Statistics:")
    print(tn_budgets.describe())

except Exception as e:
    print(f"Error loading or processing the file: {e}")
```

The Numbers Data Preview:

	id	release_date	movie	production_budget	domestic_gross	worldwide_gross
0	1	Dec 18, 2009	Avatar	\$425,000,000	\$760,507,625	\$2,776,345,279
1	2	May 20, 2011	Pirates of the Caribbean: On Stranger Tides	\$410,600,000	\$241,063,875	\$1,045,663,875
2	3	Jun 7, 2019	Dark Phoenix	\$350,000,000	\$42,762,350	\$149,762,350
3	4	May 1, 2015	Avengers: Age of Ultron	\$330,600,000	\$459,005,868	\$1,403,013,963
4	5	Dec 15, 2017	Star Wars Ep. VIII: The Last Jedi	\$317,000,000	\$620,181,382	\$1,316,721,747

```
Dataset Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5782 entries, 0 to 5781
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    5782 non-null   int64
1   release_date          5782 non-null   object
2   movie                 5782 non-null   object
3   production_budget     5782 non-null   object
4   domestic_gross        5782 non-null   object
5   worldwide_gross       5782 non-null   object
dtypes: int64(1), object(5)
memory usage: 271.2+ KB
```

Number of duplicate rows: 0

```
Missing values per column:
id                0
release_date      0
movie             0
production_budget 0
domestic_gross    0
worldwide_gross   0
dtype: int64
```

```
Summary Statistics:
           id
count  5782.000000
mean    50.372363
std     28.821076
min      1.000000
25%     25.000000
50%     50.000000
75%     75.000000
max     100.000000
```

2.3. Rotten Tomatoes Movie Info Data

```
In [4]: # Load Rotten Tomatoes Reviews dataset
rt_reviews = pd.read_csv('C:/Users/USER/Desktop/Movie-Project/data/raw/zippedData/r
print("Rotten Tomatoes Reviews Data:")
display(rt_reviews.head(), "\n") # Display the first few rows
print(rt_reviews.info()) # Get an overview of the dataset
```

Rotten Tomatoes Reviews Data:

	id	review	rating	fresh	critic	top_critic	publisher	date
0	3	A distinctly gallows take on contemporary fina...	3/5	fresh	PJ Nabarro	0	Patrick Nabarro	November 10, 2018
1	3	It's an allegory in search of a meaning that n...	NaN	rotten	Annalee Newitz	0	io9.com	May 23, 2018
2	3	... life lived in a bubble in financial dealin...	NaN	fresh	Sean Axmaker	0	Stream on Demand	January 4, 2018
3	3	Continuing along a line introduced in last yea...	NaN	fresh	Daniel Kasman	0	MUBI	November 16, 2017
4	3	... a perverse twist on neorealism...	NaN	fresh	NaN	0	Cinema Scope	October 12, 2017

'\n'

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 54432 entries, 0 to 54431

Data columns (total 8 columns):

#	Column	Non-Null Count	Dtype
0	id	54432 non-null	int64
1	review	48869 non-null	object
2	rating	40915 non-null	object
3	fresh	54432 non-null	object
4	critic	51710 non-null	object
5	top_critic	54432 non-null	int64
6	publisher	54123 non-null	object
7	date	54432 non-null	object

dtypes: int64(2), object(6)

memory usage: 3.3+ MB

None

2.4. TheMovieDB Data

```
In [5]: # Load Rotten Tomatoes Movie Info dataset
rt_info = pd.read_csv('C:/Users/USER/Desktop/Movie-Project/data/raw/zippedData/rt.m
print("Rotten Tomatoes Movie Info Data:")
display(rt_info.head()) # Display the first few rows
print(rt_info.info()) # Get an overview of the dataset
```

Rotten Tomatoes Movie Info Data:

	id	synopsis	rating	genre	director	writer	theater_date	dvd_date	curi
0	1	This gritty, fast-paced, and innovative police...	R	Action and Adventure Classics Drama	William Friedkin	Ernest Tidyman	Oct 9, 1971	Sep 25, 2001	
1	3	New York City, not-too-distant-future: Eric Pa...	R	Drama Science Fiction and Fantasy	David Cronenberg	David Cronenberg Don DeLillo	Aug 17, 2012	Jan 1, 2013	
2	5	Illeana Douglas delivers a superb performance ...	R	Drama Musical and Performing Arts	Allison Anders	Allison Anders	Sep 13, 1996	Apr 18, 2000	
3	6	Michael Douglas runs afoul of a treacherous su...	R	Drama Mystery and Suspense	Barry Levinson	Paul Attanasio Michael Crichton	Dec 9, 1994	Aug 27, 1997	
4	7	NaN	NR	Drama Romance	Rodney Bennett	Giles Cooper	NaN	NaN	

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1560 entries, 0 to 1559
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   id               1560 non-null   int64
1   synopsis         1498 non-null   object
2   rating           1557 non-null   object
3   genre            1552 non-null   object
4   director         1361 non-null   object
5   writer           1111 non-null   object
6   theater_date     1201 non-null   object
7   dvd_date         1201 non-null   object
8   currency         340 non-null    object
9   box_office       340 non-null    object
10  runtime          1530 non-null   object
11  studio           494 non-null    object
dtypes: int64(1), object(11)
memory usage: 146.4+ KB
None
```

2.5. tmdb.movies

```
In [6]: # Load TMDB dataset
```

```
tmdb_movies = pd.read_csv('C:/Users/USER/Desktop/Movie-Project/data/raw/zippedData/')
print("TheMovieDB Data:")
print(tmdb_movies.info()) # Get an overview of the dataset
display(tmdb_movies.head(), "\n") # Display the first few rows
```

TheMovieDB Data:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 26517 entries, 0 to 26516

Data columns (total 10 columns):

#	Column	Non-Null Count	Dtype
0	Unnamed: 0	26517 non-null	int64
1	genre_ids	26517 non-null	object
2	id	26517 non-null	int64
3	original_language	26517 non-null	object
4	original_title	26517 non-null	object
5	popularity	26517 non-null	float64
6	release_date	26517 non-null	object
7	title	26517 non-null	object
8	vote_average	26517 non-null	float64
9	vote_count	26517 non-null	int64

dtypes: float64(2), int64(3), object(5)

memory usage: 2.0+ MB

None

	Unnamed: 0	genre_ids	id	original_language	original_title	popularity	release_date	
0	0	[12, 14, 10751]	12444	en	Harry Potter and the Deathly Hallows: Part 1	33.533	2010-11-19	h
1	1	[14, 12, 16, 10751]	10191	en	How to Train Your Dragon	28.734	2010-03-26	
2	2	[12, 28, 878]	10138	en	Iron Man 2	28.515	2010-05-07	Ir
3	3	[16, 35, 10751]	862	en	Toy Story	28.005	1995-11-22	Tc
4	4	[28, 878, 12]	27205	en	Inception	27.920	2010-07-16	In

'\n'

2.6. im.db.zip

- Zipped SQLite database

In [7]: # Path to the SQL database file


```
db_path = 'C:/Users/USER/Desktop/dsc-phase-2-project-v3-main/unzipped/im.db'

try:
    # Connecting to the database
    conn = sqlite3.connect(db_path)
    print("Database connection successful.")

    # Check available tables in the database
    tables = pd.read_sql_query("SELECT name FROM sqlite_master WHERE type='table';")
    print("Available tables in the database:")
    print(tables)

    # Load specific tables
    movie_basics = pd.read_sql_query("SELECT * FROM movie_basics", conn)
    movie_ratings = pd.read_sql_query("SELECT * FROM movie_ratings", conn)

except sqlite3.Error as e:
    print(f"Database connection error: {e}")

finally:
    # Close the connection
    if 'conn' in locals() and conn:
        conn.close()
    print("Database connection closed.")
```

Database connection successful.

Available tables in the database:

	name
0	movie_basics
1	directors
2	known_for
3	movie_akas
4	movie_ratings
5	persons
6	principals
7	writers

Database connection closed.

```
In [8]: # Preview the `movie_basic` data
print("\nMovie Basics Preview:")
display(movie_basics.head())
movie_basics.info()
```

Movie Basics Preview:

	movie_id	primary_title	original_title	start_year	runtime_minutes	genres
0	tt0063540	Sunghursh	Sunghursh	2013	175.0	Action,Crime,Drama
1	tt0066787	One Day Before the Rainy Season	Ashad Ka Ek Din	2019	114.0	Biography,Drama
2	tt0069049	The Other Side of the Wind	The Other Side of the Wind	2018	122.0	Drama
3	tt0069204	Sabse Bada Sukh	Sabse Bada Sukh	2018	NaN	Comedy,Drama
4	tt0100275	The Wandering Soap Opera	La Telenovela Errante	2017	80.0	Comedy,Drama,Fantasy

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 146144 entries, 0 to 146143
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  -
0   movie_id              146144 non-null object
1   primary_title         146144 non-null object
2   original_title        146123 non-null object
3   start_year            146144 non-null int64
4   runtime_minutes       114405 non-null float64
5   genres                140736 non-null object
dtypes: float64(1), int64(1), object(4)
memory usage: 6.7+ MB
```

```
In [9]: # Preview the `movie_ratings` data
print("\nMovie Ratings Preview:")
display(movie_ratings.head())# Display the first few rows
movie_ratings.info()# Get an overview of the dataset
```

Movie Ratings Preview:

	movie_id	averagerating	numvotes
0	tt10356526	8.3	31
1	tt10384606	8.9	559
2	tt1042974	6.4	20
3	tt1043726	4.2	50352
4	tt1060240	6.5	21

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 73856 entries, 0 to 73855
Data columns (total 3 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   movie_id        73856 non-null  object
 1   averagerating   73856 non-null  float64
 2   numvotes        73856 non-null  int64
dtypes: float64(1), int64(1), object(1)
memory usage: 1.7+ MB
```



Since I have all the data loaded I shall proceed to cleaning it throughly.

Data Cleaning

3.1 Box Office Mojo Data

```
In [10]: # Convert 'foreign_gross' to numeric by removing non-numeric characters
bom_gross['foreign_gross'] = bom_gross['foreign_gross'].replace('[^0-9]', '', regex

# Convert 'year' column to integer type, handling non-convertible values
bom_gross['year'] = pd.to_numeric(bom_gross['year'], errors='coerce').astype('Int64

#Remove special characters and multiple spaces to standardize title
bom_gross['title'] = (
    bom_gross['title']
    .str.lower()
    .str.strip()
```

```
.str.replace(r'^\w\s', '', regex=True)
.str.replace(r'\s+', ' ', regex=True)
)

# Drop duplicate rows
bom_gross = bom_gross.drop_duplicates()

# Remove all rows with any NaN values
bom_gross.dropna(inplace=True)

# Display the first few rows of the data
display(bom_gross.head())

# Check for missing values
print(bom_gross.isnull().sum())

# Display dataset info
bom_gross.info()

# Save the DataFrame to a CSV file
bom_gross.to_csv('C:/Users/USER/Desktop/movie_insights/zippedData/processed/bom_gro
print("Dataset information saved as 'bom_gross_info.csv'.")
```

	title	studio	domestic_gross	foreign_gross	year
0	toy story 3	BV	415000000.0	652000000.0	2010
1	alice in wonderland 2010	BV	334200000.0	691300000.0	2010
2	harry potter and the deathly hallows part 1	WB	296000000.0	664300000.0	2010
3	inception	WB	292600000.0	535700000.0	2010
4	shrek forever after	P/DW	238700000.0	513900000.0	2010

```
title      0
studio     0
domestic_gross  0
foreign_gross  0
year       0
```

```
dtype: int64
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 2007 entries, 0 to 3353
```

```
Data columns (total 5 columns):
```

#	Column	Non-Null Count	Dtype
0	title	2007 non-null	object
1	studio	2007 non-null	object
2	domestic_gross	2007 non-null	float64
3	foreign_gross	2007 non-null	float64
4	year	2007 non-null	Int64

```
dtypes: Int64(1), float64(2), object(2)
```

```
memory usage: 96.0+ KB
```

```
Dataset information saved as 'bom_gross_info.csv'.
```

3.2. The Numbers Data

```
In [11]: # Remove '$' and ',' from financial columns and convert them to numeric
for col in ['production_budget', 'domestic_gross', 'worldwide_gross']:
    tn_budgets[col] = tn_budgets[col].replace(r'[\$,]', '', regex=True).astype(float)

# Convert 'release_date' to datetime
tn_budgets['release_date'] = pd.to_datetime(tn_budgets['release_date'], errors='coerce')

# Extract year from release_date
tn_budgets['year'] = tn_budgets['release_date'].dt.year

# Remove special characters and multiple spaces to standardize 'movie'
tn_budgets['movie'] = (
    tn_budgets['movie']
    .str.lower()
    .str.strip()
    .str.replace(r'^\w\s', '', regex=True)
    .str.replace(r'\s+', ' ', regex=True)
)

# Drop rows with missing financial data or release_date
tn_budgets = tn_budgets.dropna(subset=['production_budget', 'domestic_gross', 'worldwide_gross'])

# Display the first few rows of the data
display(tn_budgets.head())

# Check for missing values
print("\nMissing values per column:")
print(tn_budgets.isnull().sum())

# Display dataset info
print("\nDataset Info:")
tn_budgets.info()

# Summarize financial columns
print("\nFinancial Columns Summary:")
display(tn_budgets[['production_budget', 'domestic_gross', 'worldwide_gross']].describe())

# Save the DataFrame to a CSV file
tn_budgets.to_csv('C:/Users/USER/Desktop/movie_insights/zippedData/processed/tn_budgets.csv')
print("Dataset saved as 'tn_budgets.csv'.")
```

	id	release_date	movie	production_budget	domestic_gross	worldwide_gross	year
0	1	2009-12-18	avatar	425000000.0	760507625.0	2.776345e+09	2009
1	2	2011-05-20	pirates of the caribbean on stranger tides	410600000.0	241063875.0	1.045664e+09	2011
2	3	2019-06-07	dark phoenix	350000000.0	42762350.0	1.497624e+08	2019
3	4	2015-05-01	avengers age of ultron	330600000.0	459005868.0	1.403014e+09	2015
4	5	2017-12-15	star wars ep viii the last jedi	317000000.0	620181382.0	1.316722e+09	2017

Missing values per column:
id 0
release_date 0
movie 0
production_budget 0
domestic_gross 0
worldwide_gross 0
year 0
dtype: int64

Dataset Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5782 entries, 0 to 5781
Data columns (total 7 columns):
Column Non-Null Count Dtype

0 id 5782 non-null int64
1 release_date 5782 non-null datetime64[ns]
2 movie 5782 non-null object
3 production_budget 5782 non-null float64
4 domestic_gross 5782 non-null float64
5 worldwide_gross 5782 non-null float64
6 year 5782 non-null int32
dtypes: datetime64[ns](1), float64(3), int32(1), int64(1), object(1)
memory usage: 293.7+ KB

Financial Columns Summary:

	production_budget	domestic_gross	worldwide_gross
count	5.782000e+03	5.782000e+03	5.782000e+03
mean	3.158776e+07	4.187333e+07	9.148746e+07
std	4.181208e+07	6.824060e+07	1.747200e+08
min	1.100000e+03	0.000000e+00	0.000000e+00
25%	5.000000e+06	1.429534e+06	4.125415e+06
50%	1.700000e+07	1.722594e+07	2.798445e+07
75%	4.000000e+07	5.234866e+07	9.764584e+07
max	4.250000e+08	9.366622e+08	2.776345e+09

Dataset saved as 'tn_budgets.csv'.

3.3. TMDb Dataset

```
In [12]: # Drop the unnecessary 'Unnamed: 0' & 'genre_ids' columns
tmdb_movies.drop(columns=['Unnamed: 0', 'genre_ids'], inplace=True)

# Convert 'release_date' to datetime
tmdb_movies['release_date'] = pd.to_datetime(tmdb_movies['release_date'], errors='c

# Remove rows with invalid release dates
tmdb_movies = tmdb_movies.dropna(subset=['release_date'])

# Remove special characters and multiple spaces to standardize 'title'
tmdb_movies['title'] = (
    tmdb_movies['title']
    .str.lower()
    .str.strip()
    .str.replace(r'^\w\s', '', regex=True)
    .str.replace(r'\s+', ' ', regex=True)
)

# Drop rows with missing titles
tmdb_movies = tmdb_movies.dropna(subset=['title'])

# Display the first few rows of the cleaned data
display(tmdb_movies.head())

# Check for missing values in the dataset
print("\nMissing values per column:")
print(tmdb_movies.isnull().sum())

# Display dataset information
print("\nDataset Info:")
tmdb_movies.info()

# Summarize numerical and categorical columns
print("\nSummary Statistics:")
```

```
display(tmdb_movies.describe(include='all'))

# Save the DataFrame to a CSV file
tmdb_movies.to_csv('C:/Users/USER/Desktop/movie_insights/zippedData/processed/tmdb_
print("Dataset saved as 'tmdb_movies.csv'.")
```

	id	original_language	original_title	popularity	release_date	title	vote_average
0	12444	en	Harry Potter and the Deathly Hallows: Part 1	33.533	2010-11-19	harry potter and the deathly hallows part 1	7.7
1	10191	en	How to Train Your Dragon	28.734	2010-03-26	how to train your dragon	7.7
2	10138	en	Iron Man 2	28.515	2010-05-07	iron man 2	6.8
3	862	en	Toy Story	28.005	1995-11-22	toy story	7.9
4	27205	en	Inception	27.920	2010-07-16	inception	8.3

Missing values per column:

id	0
original_language	0
original_title	0
popularity	0
release_date	0
title	0
vote_average	0
vote_count	0
dtype:	int64

Dataset Info:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 26517 entries, 0 to 26516
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     26517 non-null  int64
1   original_language      26517 non-null  object
2   original_title         26517 non-null  object
3   popularity             26517 non-null  float64
4   release_date           26517 non-null  datetime64[ns]
5   title                  26517 non-null  object
6   vote_average           26517 non-null  float64
7   vote_count             26517 non-null  int64
dtypes: datetime64[ns](1), float64(2), int64(2), object(3)
memory usage: 1.6+ MB
```

Summary Statistics:

	id	original_language	original_title	popularity	release_date	ti
count	26517.000000	26517	26517	26517.000000	26517	26517
unique	NaN	76	24835	NaN	NaN	24835
top	NaN	en	Eden	NaN	NaN	ho
freq	NaN	23291	7	NaN	NaN	
mean	295050.153260	NaN	NaN	3.130912	2014-06-10 02:50:14.730173184	N
min	27.000000	NaN	NaN	0.600000	1930-04-29 00:00:00	N
25%	157851.000000	NaN	NaN	0.600000	2012-06-29 00:00:00	N
50%	309581.000000	NaN	NaN	1.374000	2014-09-19 00:00:00	N
75%	419542.000000	NaN	NaN	3.694000	2016-10-01 00:00:00	N
max	608444.000000	NaN	NaN	80.773000	2020-12-25 00:00:00	N
std	153661.615648	NaN	NaN	4.355229	NaN	N

Dataset saved as 'tmdb_movies.csv'.

3.4. Rotten Tomatoes Reviews Data

```
In [13]: # Drop rows where `review`, `rating`, or `critic` is missing
rt_reviews.dropna(subset=['review', 'rating', 'critic', 'publisher'], inplace=True)

# Convert `date` column to datetime
rt_reviews['date'] = pd.to_datetime(rt_reviews['date'], errors='coerce')

# Parse `rating` to extract numeric scores (e.g., '3/5' -> 3.0)
def parse_rating(rating):
    try:
        return float(rating.split('/')[0]) if '/' in rating else None
    except:
        return None

rt_reviews['rating'] = rt_reviews['rating'].apply(parse_rating)

# Drop rows with missing `rating_score`
rt_reviews.dropna(subset=['rating'], inplace=True)

# Normalize ratings to a 0-10 scale
def normalize_rating(rating):
    try:
        parts = rating.split('/')
        if len(parts) == 2 and float(parts[1]) > 0:
```

```
        return (float(parts[0]) / float(parts[1])) * 10
    except:
        return None
    return None

rt_reviews['normalized_rating'] = rt_reviews['rating'].apply(normalize_rating)

# Drop rows with invalid or missing dates
rt_reviews = rt_reviews.dropna(subset=['date'])

# Remove duplicates
rt_reviews.drop_duplicates(inplace=True)

# Rename Columns to snake_case
rt_reviews.rename(columns={
    'review': 'review_text',
    'rating': 'rating_score',
    'fresh': 'is_fresh',
    'critic': 'critic_name',
    'top_critic': 'is_top_critic',
    'publisher': 'publisher_name',
    'date': 'review_date'
}, inplace=True)

# Display dataset info
print("Rotten Tomatoes Reviews Data:")
print(rt_reviews.info())

# Display summary statistics for numeric columns
print("\nSummary Statistics:")
print(rt_reviews[['rating_score', 'normalized_rating']].describe())

# Display the first few rows of the cleaned data
display(rt_reviews.head())

# Save the DataFrame to a CSV file
output_path = 'C:/Users/USER/Desktop/movie_insights/zippedData/processed/rt_reviews'
rt_reviews.to_csv(output_path, index=False)
print(f"Cleaned dataset saved as '{output_path}'. Total rows: {rt_reviews.shape[0]}")

# Save the DataFrame to a CSV file
rt_reviews.to_csv('C:/Users/USER/Desktop/movie_insights/zippedData/processed/rt_rev
print("Dataset saved as 'rt_reviews.csv'.")
```

Rotten Tomatoes Reviews Data:
<class 'pandas.core.frame.DataFrame'>
Index: 27587 entries, 0 to 54424
Data columns (total 9 columns):

#	Column	Non-Null Count	Dtype
0	id	27587 non-null	int64
1	review_text	27587 non-null	object
2	rating_score	27587 non-null	float64
3	is_fresh	27587 non-null	object
4	critic_name	27587 non-null	object
5	is_top_critic	27587 non-null	int64
6	publisher_name	27587 non-null	object
7	review_date	27587 non-null	datetime64[ns]
8	normalized_rating	0 non-null	object

dtypes: datetime64[ns](1), float64(1), int64(2), object(5)
memory usage: 2.1+ MB
None

Summary Statistics:

	rating_score
count	27587.000000
mean	3.140461
std	1.505725
min	0.000000
25%	2.000000
50%	3.000000
75%	4.000000
max	9.600000

	id	review_text	rating_score	is_fresh	critic_name	is_top_critic	publisher_name	review
0	3	A distinctly gallows take on contemporary fina...	3.0	fresh	PJ Nabarro	0	Patrick Nabarro	2018
7	3	Cronenberg is not a director to be daunted by ...	2.0	rotten	Matt Kelemen	0	Las Vegas CityLife	2013
12	3	Robert Pattinson works mighty hard to make Cos...	2.0	rotten	Christian Toto	0	Big Hollywood	2013
14	3	For those who like their Cronenberg thick and ...	3.0	fresh	Marty Mapes	0	Movie Habit	2012
15	3	For better or worse - often both - Cosmopolis ...	3.0	fresh	Adam Ross	0	The Aristocrat	2012

Cleaned dataset saved as 'C:/Users/USER/Desktop/movie_insights/zippedData/processed/rt_reviews.csv'. Total rows: 27587
Dataset saved as 'rt_reviews.csv'.



3.5. Rotten Tomatoes Movie Info Data

```
In [14]: # Convert to datetime and coerce invalid dates to NaT
rt_info['theater_date'] = pd.to_datetime(rt_info['theater_date'], errors='coerce')
rt_info['dvd_date'] = pd.to_datetime(rt_info['dvd_date'], errors='coerce')

# Clean 'runtime' to extract numerical values
rt_info['runtime'] = rt_info['runtime'].str.extract(r'(\d+)').astype(float) # Use

# Clean 'box_office' to extract numerical values
rt_info['box_office'] = rt_info['box_office'].replace(r'[\$,]', '', regex=True).ast

# Drop all rows with any NaN values
rt_info.dropna(inplace=True)

# Display the cleaned dataset
print("After dropping all rows with NaN values:")
print(rt_info.info())
display(rt_info.head())

# Save the DataFrame to a CSV file
rt_reviews.to_csv('C:/Users/USER/Desktop/movie_insights/zippedData/processed/rt_inf
print("Dataset saved as 'rt_info.csv'.")
```

```
After dropping all rows with NaN values:
<class 'pandas.core.frame.DataFrame'>
Index: 235 entries, 1 to 1545
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   id               235 non-null   int64
1   synopsis         235 non-null   object
2   rating           235 non-null   object
3   genre            235 non-null   object
4   director         235 non-null   object
5   writer           235 non-null   object
6   theater_date     235 non-null   datetime64[ns]
7   dvd_date         235 non-null   datetime64[ns]
8   currency         235 non-null   object
9   box_office       235 non-null   float64
10  runtime          235 non-null   float64
11  studio           235 non-null   object
dtypes: datetime64[ns](2), float64(2), int64(1), object(7)
memory usage: 23.9+ KB
None
```

	id	synopsis	rating	genre	director	writer	theater_date	dvd_date	
1	3	New York City, not-too-distant-future: Eric Pa...	R	Drama Science Fiction and Fantasy	David Cronenberg	David Cronenberg Don DeLillo	2012-08-17	2013-01-01	
6	10	Some cast and crew from NBC's highly acclaimed...	PG-13	Comedy	Jake Kasdan	Mike White	2002-01-11	2002-06-18	
7	13	Stewart Kane, an Irishman living in the Austra...	R	Drama	Ray Lawrence	Raymond Carver Beatrix Christian	2006-04-27	2007-10-02	
15	22	Two-time Academy Award Winner Kevin Spacey giv...	R	Comedy Drama Mystery and Suspense	George Hickenlooper	Norman Snider	2010-12-17	2011-04-05	
18	25	From ancient Japan's most enduring tale, the e...	PG-13	Action and Adventure Drama Science Fiction and...	Carl Erik Rinsch	Chris Morgan Hossein Amini	2013-12-25	2014-04-01	

Dataset saved as 'rt_info.csv'.

3.6.im.db.zip

- Zipped SQLite database

a. movie_basics

```
In [15]: # Display the first few rows
display(movie_basics.head())
movie_basics.info()
```

	movie_id	primary_title	original_title	start_year	runtime_minutes	genres
0	tt0063540	Sunghursh	Sunghursh	2013	175.0	Action,Crime,Drama
1	tt0066787	One Day Before the Rainy Season	Ashad Ka Ek Din	2019	114.0	Biography,Drama
2	tt0069049	The Other Side of the Wind	The Other Side of the Wind	2018	122.0	Drama
3	tt0069204	Sabse Bada Sukh	Sabse Bada Sukh	2018	NaN	Comedy,Drama
4	tt0100275	The Wandering Soap Opera	La Telenovela Errante	2017	80.0	Comedy,Drama,Fantasy

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 146144 entries, 0 to 146143
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  -
0   movie_id              146144 non-null object
1   primary_title         146144 non-null object
2   original_title        146123 non-null object
3   start_year            146144 non-null int64
4   runtime_minutes       114405 non-null float64
5   genres                140736 non-null object
dtypes: float64(1), int64(1), object(4)
memory usage: 6.7+ MB
```

```
In [16]: # Drop rows where 'runtime_minutes' or 'genres' columns have missing values
movie_basics.dropna(subset=['runtime_minutes', 'genres'], inplace=True)

# Drop any remaining rows with missing values in any column
movie_basics.dropna(inplace=True)

# Normalize 'primary_title' and 'original_title' columns: Lowercase, remove punctua
string_columns = ['primary_title', 'original_title']
for col in string_columns:
    movie_basics[col] = (
        movie_basics[col]
        .str.lower() # Convert to Lowercase
        .str.strip() # Remove Leading/trailing spaces
        .str.replace(r'^\w[s]', '', regex=True) # Remove punctuation
        .str.replace(r'\s+', ' ', regex=True) # Replace multiple spaces with a sin
    )

# Display the first few rows to inspect changes
print("\nPreview of cleaned movie_basics dataset:")
display(movie_basics.head())

# Display information about the cleaned dataframe
print("\nInformation about cleaned movie_basics dataset:")
```



```
movie_basics.info()

# Save the DataFrame to a CSV file
rt_reviews.to_csv('C:/Users/USER/Desktop/movie_insights/zippedData/processed/movie_
print("Dataset saved as 'movie_basics.csv'.")
```

Preview of cleaned movie_basics dataset:

	movie_id	primary_title	original_title	start_year	runtime_minutes	genres
0	tt0063540	sunghursh	sunghursh	2013	175.0	Action,Crime,Drama
1	tt0066787	one day before the rainy season	ashad ka ek din	2019	114.0	Biography,Drama
2	tt0069049	the other side of the wind	the other side of the wind	2018	122.0	Drama
4	tt0100275	the wandering soap opera	la telenovela errante	2017	80.0	Comedy,Drama,Fantasy
5	tt0111414	a thin life	a thin life	2018	75.0	Comedy

Information about cleaned movie_basics dataset:

```
<class 'pandas.core.frame.DataFrame'>
```

Index: 112232 entries, 0 to 146139

Data columns (total 6 columns):

#	Column	Non-Null Count	Dtype
0	movie_id	112232 non-null	object
1	primary_title	112232 non-null	object
2	original_title	112232 non-null	object
3	start_year	112232 non-null	int64
4	runtime_minutes	112232 non-null	float64
5	genres	112232 non-null	object

dtypes: float64(1), int64(1), object(4)

memory usage: 6.0+ MB

Dataset saved as 'movie_basics.csv'.

b. movie_ratings

```
In [17]: # Display the first few rows of the movie_ratings dataframe
display(movie_ratings.head())

# Display detailed information about the dataframe
movie_ratings.info()

# Observed that there are no missing values in any column of the dataframe.
print(f'The movie_ratings dataframe is already clean with no missing values.')

# Save the DataFrame to a CSV file
rt_reviews.to_csv('C:/Users/USER/Desktop/movie_insights/zippedData/processed/movie_
print("Dataset saved as 'movie_ratings.csv'.")
```

	movie_id	averagerating	numvotes
0	tt10356526	8.3	31
1	tt10384606	8.9	559
2	tt1042974	6.4	20
3	tt1043726	4.2	50352
4	tt1060240	6.5	21

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 73856 entries, 0 to 73855
```

```
Data columns (total 3 columns):
```

```
#   Column          Non-Null Count  Dtype
---  -
0   movie_id        73856 non-null    object
1   averagerating    73856 non-null    float64
2   numvotes         73856 non-null    int64
```

```
dtypes: float64(1), int64(1), object(1)
```

```
memory usage: 1.7+ MB
```

```
The movie_ratings dataframe is already clean with no missing values.
```

```
Dataset saved as 'movie_ratings.csv'.
```

3.a. Merging Datasets

```
In [18]: # Prepare `bom_gross` and `tn_budgets` for merging
tn_budgets['year'] = tn_budgets['release_date'].dt.year # Extract year from release_date

# Clean the 'title' field in bom_gross
bom_gross['title'] = (
    bom_gross['title']
    .str.lower() # Convert to lowercase
    .str.strip() # Remove leading and trailing whitespace
    .str.replace(r'^\w\s', '', regex=True) # Remove special characters
    .str.replace(r'\s+', ' ', regex=True) # Replace multiple spaces with a single space
)

# Clean the 'title' field in tn_budgets
tn_budgets['movie'] = (
    tn_budgets['movie']
    .str.lower() # Convert to lowercase
    .str.strip() # Remove leading and trailing whitespace
    .str.replace(r'^\w\s', '', regex=True) # Remove special characters
    .str.replace(r'\s+', ' ', regex=True) # Replace multiple spaces with a single space
)

# Check Overlap Between bom_gross and tmdb_movies:
matched_titles = bom_gross['title'].isin(tmdb_movies['title']).sum()
print(f"Number of matched titles between bom_gross and tmdb_movies: {matched_titles}")
```

```
Number of matched titles between bom_gross and tmdb_movies: 1560
```

```
In [19]: # Check Overlap Between bom_gross and tmdb_movies:
matched_titles = movie_basics['original_title'].isin(tmdb_movies['original_title'])
print(f"Number of matched titles between bom_gross and tmdb_movies: {matched_titles}")
```

Number of matched titles between bom_gross and tmdb_movies: 67

```
In [20]: #Check Overlap Between tn_budgets and tmdb_movies:
matched_titles = tn_budgets['movie'].isin(tmdb_movies['title']).sum()
print(f"Number of matched titles between tn_budgets and tmdb_movies: {matched_title
```

Number of matched titles between tn_budgets and tmdb_movies: 2051

```
In [21]: #Check for duplicates in title_normalized:

display("Duplicates in bom_gross:", bom_gross[bom_gross['title'].duplicated()])
display("Duplicates in tn_budgets:", tn_budgets[tn_budgets['movie'].duplicated()])
display("Duplicates in tmdb_movies:", tmdb_movies[tmdb_movies['title'].duplicated()])
```

'Duplicates in bom_gross:'

title	studio	domestic_gross	foreign_gross	year			
'Duplicates in tn_budgets:'							
	id	release_date	movie	production_budget	domestic_gross	worldwide_gross	year
273	74	1998-05-19	godzilla	125000000.0	136314294.0	376000000.0	1998
408	9	2018-11-21	robin hood	99000000.0	30824628.0	84747441.0	2018
484	85	2005-07-08	fantastic four	87500000.0	154696080.0	333132750.0	2005
543	44	1999-05-07	the mummy	80000000.0	155385488.0	416385488.0	1999
707	8	1997-06-13	hercules	70000000.0	99112101.0	250700000.0	1997
...
5668	69	1942-11-16	cat people	134000.0	4000000.0	8000000.0	1942
5676	77	1968-10-01	night of the living dead	114000.0	12087064.0	30087064.0	1968
5677	78	1915-02-08	the birth of a nation	110000.0	10000000.0	11000000.0	1915
5699	100	1972-08-30	the last house on the left	87000.0	3100000.0	3100000.0	1972
5718	19	2008-02-22	the signal	50000.0	251150.0	406299.0	2008

84 rows × 7 columns

'Duplicates in tmdb_movies:'

	id	original_language	original_title	popularity	release_date	title	vote_ave
781	51462	en	Brotherhood	2.235	2010-01-03	brotherhood	
1037	44369	tl	Boy	1.504	2009-06-01	boy	
1230	371702	en	All That Glitters	1.241	2010-09-25	all that glitters	
1354	155711	en	After-Life	0.994	2010-01-01	afterlife	
1501	36410	en	Zero	0.840	2010-02-06	zero	
...	
26495	556601	en	Recursion	0.600	2018-08-28	recursion	
26504	534282	en	Head	0.600	2015-03-28	head	
26506	561861	en	Eden	0.600	2018-11-25	eden	
26510	495045	en	Fail State	0.600	2018-10-19	fail state	
26511	492837	en	Making Filmmakers	0.600	2018-04-07	making filmmakers	

1880 rows × 8 columns

```
In [22]: #Remove duplicates:
tn_budgets.drop_duplicates(subset='movie', inplace=True)
tmdb_movies.drop_duplicates(subset='title', inplace=True)
```

```
In [23]: #Check for duplicates in title_normalized:

display("Duplicates in bom_gross:", bom_gross[bom_gross['title'].duplicated()])
display("Duplicates in tn_budgets:", tn_budgets[tn_budgets['movie'].duplicated()])
display("Duplicates in tmdb_movies:", tmdb_movies[tmdb_movies['title'].duplicated()])
```

'Duplicates in bom_gross:'

title	studio	domestic_gross	foreign_gross	year
-------	--------	----------------	---------------	------

'Duplicates in tn_budgets:'

id	release_date	movie	production_budget	domestic_gross	worldwide_gross	year
----	--------------	-------	-------------------	----------------	-----------------	------

'Duplicates in tmdb_movies:'

id	original_language	original_title	popularity	release_date	title	vote_average	vote_count
----	-------------------	----------------	------------	--------------	-------	--------------	------------

Join bom_gross and tmdb_movies

```
In [24]: #Join bom_gross and tmdb_movies
bom_tmdb_merged = pd.merge(bom_gross, tmdb_movies, on='title', how='inner', suffixes=('_bom', '_tmdb'))
print(f"bom_tmdb_merged shape: {bom_tmdb_merged.shape}")

# Check for missing values
```

```
missing_values = bom_tmdb_merged.isna().sum()
print("\nMissing values in bom_tmdb_merged:")
print(missing_values)

# Preview merged data
print("\nPreview of merged dataset:")
display(bom_tmdb_merged.head())
```

bom_tmdb_merged shape: (1560, 12)

Missing values in bom_tmdb_merged:

```
title          0
studio          0
domestic_gross 0
foreign_gross  0
year           0
id             0
original_language 0
original_title  0
popularity     0
release_date   0
vote_average   0
vote_count     0
dtype: int64
```

Preview of merged dataset:

	title	studio	domestic_gross	foreign_gross	year	id	original_language	original_
0	toy story 3	BV	415000000.0	652000000.0	2010	10193	en	Toy Stc
1	harry potter and the deathly hallows part 1	WB	296000000.0	664300000.0	2010	12444	en	Harry Po anc Dei Hallows:
2	inception	WB	292600000.0	535700000.0	2010	27205	en	Incep
3	shrek forever after	P/DW	238700000.0	513900000.0	2010	10192	en	S Forever /
4	the twilight saga eclipse	Sum.	300500000.0	398000000.0	2010	24021	en	The Twi Saga: Ec

'movie' from tn_budgets and 'title' from tmdb_movies

```
In [25]: # Merge using 'movie' from tn_budgets and 'title' from tmdb_movies
budgets_tmdb_merged = pd.merge(
    tn_budgets, tmdb_movies,
    left_on='movie', right_on='title',
```

```
        how='inner', suffixes=('_budget', '_tmdb')
    )

    # Print shape of the merged dataframe
    print(f"budgets_tmdb_merged shape: {budgets_tmdb_merged.shape}")
```

budgets_tmdb_merged shape: (1998, 15)

```
In [26]: # Check for missing values in the merged dataset
missing_values = budgets_tmdb_merged.isna().sum()
print("\nMissing values in budgets_tmdb_merged:")
print(missing_values)

#Preview merged dataset
print("\nPreview of merged dataset:")
display(budgets_tmdb_merged.head())
```

Missing values in budgets_tmdb_merged:

```
id_budget      0
release_date_budget  0
movie          0
production_budget  0
domestic_gross  0
worldwide_gross  0
year           0
id_tmdb        0
original_language  0
original_title  0
popularity     0
release_date_tmdb  0
title          0
vote_average   0
vote_count     0
dtype: int64
```

Preview of merged dataset:

	id_budget	release_date_budget	movie	production_budget	domestic_gross	worldwide
0	1	2009-12-18	avatar	425000000.0	760507625.0	2.77634
1	2	2011-05-20	pirates of the caribbean on stranger tides	410600000.0	241063875.0	1.04566
2	4	2015-05-01	avengers age of ultron	330600000.0	459005868.0	1.40301
3	7	2018-04-27	avengers infinity war	300000000.0	678815482.0	2.04813
4	9	2017-11-17	justice league	300000000.0	229024295.0	6.55945

```
In [27]: # Combine all three datasets
final_merged = pd.merge(
    bom_tmdb_merged, budgets_tmdb_merged,
    on='title',
    how='inner',
```

```
    suffixes=('_bom_tmdb', '_budget_tmdb')
)
print(f"final_merged shape: {final_merged.shape}")
```

final_merged shape: (1082, 26)

```
In [28]: #Check for missing values
print("\nMissing values in final_merged:")
print(final_merged.isna().sum())

# Preview the final merged dataset
display(final_merged.head())
```

Missing values in final_merged:

title	0
studio	0
domestic_gross_bom_tmdb	0
foreign_gross	0
year_bom_tmdb	0
id	0
original_language_bom_tmdb	0
original_title_bom_tmdb	0
popularity_bom_tmdb	0
release_date	0
vote_average_bom_tmdb	0
vote_count_bom_tmdb	0
id_budget	0
release_date_budget	0
movie	0
production_budget	0
domestic_gross_budget_tmdb	0
worldwide_gross	0
year_budget_tmdb	0
id_tmdb	0
original_language_budget_tmdb	0
original_title_budget_tmdb	0
popularity_budget_tmdb	0
release_date_tmdb	0
vote_average_budget_tmdb	0
vote_count_budget_tmdb	0

dtype: int64

	title	studio	domestic_gross_bom_tmdb	foreign_gross	year_bom_tmdb	id	origir
0	toy story 3	BV	415000000.0	652000000.0	2010	10193	
1	inception	WB	292600000.0	535700000.0	2010	27205	
2	shrek forever after	P/DW	238700000.0	513900000.0	2010	10192	
3	the twilight saga eclipse	Sum.	300500000.0	398000000.0	2010	24021	
4	iron man 2	Par.	312400000.0	311500000.0	2010	10138	

5 rows × 26 columns

```
In [29]: # Remove unnecessary or duplicate columns
columns_to_drop = [
    'id', 'id_budget', 'id_tmdb', 'release_date_tmdb', 'year_budget_tmdb', 'domestic_
    'original_language_budget_tmdb', 'original_title_bom_tmdb', 'vote_count_budget_
    'release_date_budget', 'original_title_budget_tmdb'
]

final_merged_cleaned = final_merged.drop(columns=columns_to_drop)

# Rename columns for consistency and clarity
final_merged_cleaned.rename(columns={
    'release_date': 'release_date',
    'year_bom_tmdb': 'year',
    'vote_count_bom_tmdb': 'vote_count',
    'vote_average_budget_tmdb': 'vote_average',
    'domestic_gross_budget_tmdb': 'domestic_gross',
    'original_language_bom_tmdb': 'original_language',
    'popularity_bom_tmdb': 'popularity'
}, inplace=True)

# Display the cleaned dataframe information
print("Cleaned final_merged DataFrame Info:")
final_merged_cleaned.info()

# Preview the cleaned dataframe
display(final_merged_cleaned.head())
```

Cleaned final_merged DataFrame Info:

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 1082 entries, 0 to 1081

Data columns (total 13 columns):

#	Column	Non-Null Count	Dtype
0	title	1082 non-null	object
1	studio	1082 non-null	object
2	foreign_gross	1082 non-null	float64
3	year	1082 non-null	Int64
4	original_language	1082 non-null	object
5	popularity	1082 non-null	float64
6	release_date	1082 non-null	datetime64[ns]
7	vote_count	1082 non-null	int64
8	movie	1082 non-null	object
9	production_budget	1082 non-null	float64
10	domestic_gross	1082 non-null	float64
11	worldwide_gross	1082 non-null	float64
12	vote_average	1082 non-null	float64

dtypes: Int64(1), datetime64[ns](1), float64(6), int64(1), object(4)

memory usage: 111.1+ KB

	title	studio	foreign_gross	year	original_language	popularity	release_date	vote_co
0	toy story 3	BV	652000000.0	2010	en	24.445	2010-06-17	8
1	inception	WB	535700000.0	2010	en	27.920	2010-07-16	26
2	shrek forever after	P/DW	513900000.0	2010	en	15.041	2010-05-16	3
3	the twilight saga eclipse	Sum.	398000000.0	2010	en	20.340	2010-06-23	4
4	iron man 2	Par.	311500000.0	2010	en	28.515	2010-05-07	16

```
In [30]: # Save the DataFrame to a CSV file
final_merged_cleaned.to_csv('C:/Users/USER/Desktop/movie_insights/zippedData/processedData/final_merged_cleaned.csv')
print("Dataset saved as 'final_merged_cleaned.csv'.")
```

Dataset saved as 'final_merged_cleaned.csv'.

```
In [31]: #Check Overlap Between movie_basic and movie_ratings:
matched_titles = movie_basics['movie_id'].isin(movie_ratings['movie_id']).sum()
print(f"Number of matched movie_id between movie_basics and movie_ratings: {matched_titles}")
```

Number of matched movie_id between movie_basics and movie_ratings: 65720

```
In [32]: # Merge `movie_basics` and `movie_ratings` on `movie_id`
db_data = pd.merge(movie_basics, movie_ratings, on='movie_id', how='inner')

# Display the merged DataFrame information
```

```
print("Merged db_data DataFrame Info:")
db_data.info()

# Display the first few rows of the merged DataFrame
display(db_data.head())
```

Merged db_data DataFrame Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 65720 entries, 0 to 65719

Data columns (total 8 columns):

#	Column	Non-Null Count	Dtype
0	movie_id	65720 non-null	object
1	primary_title	65720 non-null	object
2	original_title	65720 non-null	object
3	start_year	65720 non-null	int64
4	runtime_minutes	65720 non-null	float64
5	genres	65720 non-null	object
6	averagerating	65720 non-null	float64
7	numvotes	65720 non-null	int64

dtypes: float64(2), int64(2), object(4)

memory usage: 4.0+ MB

	movie_id	primary_title	original_title	start_year	runtime_minutes	genres
0	tt0063540	sunghursh	sunghursh	2013	175.0	Action,Crime,D
1	tt0066787	one day before the rainy season	ashad ka ek din	2019	114.0	Biography,D
2	tt0069049	the other side of the wind	the other side of the wind	2018	122.0	D
3	tt0100275	the wandering soap opera	la telenovela errante	2017	80.0	Comedy,Drama,Fa
4	tt0137204	joe finds grace	joe finds grace	2017	83.0	Adventure,Animation,Co

```
In [33]: # Create a normalized version of the title for merging (without changing original t
db_data['primary_title'] = (
    db_data['primary_title']
    .str.lower()
    .str.strip()
    .str.replace(r'^\w\s', '', regex=True) # Remove special characters
    .str.replace(r'\s+', ' ', regex=True) # Replace multiple spaces with a single
)
```

```
In [34]: #Check Overlap Between bom_gross and tmdb_movies:
matched_titles = final_merged_cleaned['title'].isin(db_data['primary_title']).sum()
print(f"Number of matched movie_id between movie_basics and movie_ratings: {matched
```

Number of matched movie_id between movie_basics and movie_ratings: 1062

```
In [35]: # Standardize titles for merging
final_merged_cleaned['title'] = final_merged_cleaned['title'].str.lower().str.strip

# Merge datasets on titles
merged_data = pd.merge(
    final_merged_cleaned,
    db_data,
    left_on='title',
    right_on='primary_title',
    how='inner'
)
```

```
In [36]: display(merged_data.head())
```

	title	studio	foreign_gross	year	original_language	popularity	release_date	vote_co
0	toy story 3	BV	652000000.0	2010	en	24.445	2010-06-17	8
1	inception	WB	535700000.0	2010	en	27.920	2010-07-16	26
2	shrek forever after	P/DW	513900000.0	2010	en	15.041	2010-05-16	3
3	the twilight saga eclipse	Sum.	398000000.0	2010	en	20.340	2010-06-23	4
4	iron man 2	Par.	311500000.0	2010	en	28.515	2010-05-07	16

5 rows × 21 columns

```
In [37]: print(merged_data.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1241 entries, 0 to 1240
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   title                 1241 non-null   object
1   studio               1241 non-null   object
2   foreign_gross        1241 non-null   float64
3   year                 1241 non-null   Int64
4   original_language    1241 non-null   object
5   popularity           1241 non-null   float64
6   release_date         1241 non-null   datetime64[ns]
7   vote_count           1241 non-null   int64
8   movie                1241 non-null   object
9   production_budget    1241 non-null   float64
10  domestic_gross       1241 non-null   float64
11  worldwide_gross      1241 non-null   float64
12  vote_average         1241 non-null   float64
13  movie_id             1241 non-null   object
14  primary_title        1241 non-null   object
15  original_title       1241 non-null   object
16  start_year           1241 non-null   int64
17  runtime_minutes      1241 non-null   float64
18  genres               1241 non-null   object
19  averagerating        1241 non-null   float64
20  numvotes             1241 non-null   int64
dtypes: Int64(1), datetime64[ns](1), float64(8), int64(3), object(8)
memory usage: 204.9+ KB
None
```

```
In [38]: # Drop redundant columns
merged_data.drop(['original_title', 'movie', 'primary_title', 'movie_id'], axis=1, i
```

```
In [39]: merged_data.info()
merged_data.head()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1241 entries, 0 to 1240
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype
---  -
0   title                 1241 non-null   object
1   studio                1241 non-null   object
2   foreign_gross         1241 non-null   float64
3   year                  1241 non-null   Int64
4   original_language     1241 non-null   object
5   popularity            1241 non-null   float64
6   release_date          1241 non-null   datetime64[ns]
7   vote_count            1241 non-null   int64
8   production_budget     1241 non-null   float64
9   domestic_gross        1241 non-null   float64
10  worldwide_gross       1241 non-null   float64
11  vote_average          1241 non-null   float64
12  start_year            1241 non-null   int64
13  runtime_minutes       1241 non-null   float64
14  genres                1241 non-null   object
15  averagerating         1241 non-null   float64
16  numvotes              1241 non-null   int64
dtypes: Int64(1), datetime64[ns](1), float64(8), int64(3), object(4)
memory usage: 166.2+ KB
```

Out[39]:

	title	studio	foreign_gross	year	original_language	popularity	release_date	vote_
0	toy story 3	BV	652000000.0	2010	en	24.445	2010-06-17	
1	inception	WB	535700000.0	2010	en	27.920	2010-07-16	
2	shrek forever after	P/DW	513900000.0	2010	en	15.041	2010-05-16	
3	the twilight saga eclipse	Sum.	398000000.0	2010	en	20.340	2010-06-23	
4	iron man 2	Par.	311500000.0	2010	en	28.515	2010-05-07	

3.1. Derived Metrics

a. Profit and Profit Margin

```
In [40]: # Calculate profit by subtracting the production budget from the worldwide gross
merged_data['profit'] = merged_data['worldwide_gross'] - merged_data['production_bud
# 'profit' is the difference between how much the movie made globally ('worldwide_g

# Calculate profit margin by dividing the profit by the production budget
merged_data['profit_margin'] = merged_data['profit'] / merged_data['production_budg
# 'profit_margin' is the ratio of profit to production budget, indicating the retur
```

b. Release Windows

```
In [41]: # Extract the month from the 'release_date' column and store it in a new column 're
merged_data['release_month'] = merged_data['release_date'].dt.month

# Create a dictionary that maps month numbers to their corresponding month names
month_mapping = {
    1: 'January', 2: 'February', 3: 'March', 4: 'April', 5: 'May', 6: 'June',
    7: 'July', 8: 'August', 9: 'September', 10: 'October', 11: 'November', 12: 'Dec
}

# Map the 'release_month' column to month names using the dictionary
merged_data['release_month_name'] = merged_data['release_month'].map(month_mapping)

# Extract the quarter from the 'release_date' column and store it in a new column '
merged_data['release_quarter'] = merged_data['release_date'].dt.quarter
# The '.dt.quarter' function extracts the quarter of the year (1, 2, 3, or 4) based
# Quarters are typically divided as:
# - Q1: January, February, March
# - Q2: April, May, June
# - Q3: July, August, September
# - Q4: October, November, December
```

```
In [42]: merged_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1241 entries, 0 to 1240
Data columns (total 22 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   title                 1241 non-null   object 
 1   studio                1241 non-null   object 
 2   foreign_gross         1241 non-null   float64
 3   year                  1241 non-null   Int64  
 4   original_language     1241 non-null   object 
 5   popularity            1241 non-null   float64
 6   release_date          1241 non-null   datetime64[ns]
 7   vote_count            1241 non-null   int64  
 8   production_budget     1241 non-null   float64
 9   domestic_gross        1241 non-null   float64
10  worldwide_gross       1241 non-null   float64
11  vote_average          1241 non-null   float64
12  start_year            1241 non-null   int64  
13  runtime_minutes       1241 non-null   float64
14  genres                1241 non-null   object 
15  averagerating         1241 non-null   float64
16  numvotes              1241 non-null   int64  
17  profit                1241 non-null   float64
18  profit_margin         1241 non-null   float64
19  release_month         1241 non-null   int32  
20  release_month_name     1241 non-null   object 
21  release_quarter       1241 non-null   int32  
dtypes: Int64(1), datetime64[ns](1), float64(10), int32(2), int64(3), object(5)
memory usage: 204.9+ KB
```

c. Genre Analysis

```
In [44]: genres_split = merged_data.assign(genres=merged_data['genres'].str.split(',')).explode()
# Display the result
display(genres_split.head())
```

	title	studio	foreign_gross	year	original_language	popularity	release_date	vote_count
0	toy story 3	BV	652000000.0	2010	en	24.445	2010-06-17	8
0	toy story 3	BV	652000000.0	2010	en	24.445	2010-06-17	8
0	toy story 3	BV	652000000.0	2010	en	24.445	2010-06-17	8
1	inception	WB	535700000.0	2010	en	27.920	2010-07-16	2
1	inception	WB	535700000.0	2010	en	27.920	2010-07-16	2

5 rows × 22 columns

```
In [45]: # Save the DataFrame to a CSV file
merged_data.to_csv('C:/Users/USER/Desktop/movie_insights/zippedData/processed/merged_data.csv')
```



```
print("Dataset saved as 'merged_movie_clean.csv'.")
```

Dataset saved as 'merged_movie_clean.csv'.

Exploratory Data Analysis

Genre Analysis

Q1: Which genres consistently generate the highest revenue?

I will analyze which genres consistently generate the highest revenue by examining the average revenue for each genre. To identify trends, I'll use visualizations like bar plots

```
In [ ]: # Group the dataframe by 'genres' to aggregate movies by their genre
# For each genre group, I calculate the mean of 'worldwide_gross' (average revenue per movie)
genres_revenue = genres_split.groupby('genres')['worldwide_gross'].mean()

# Sort the resulting series in descending order based on the average worldwide gross
# This will display the genres with the highest average revenue at the top
genres_revenue = genres_revenue.sort_values(ascending=False)

display(genres_revenue)

# Select the top 10 genres based on average worldwide gross
top_10_genres = genres_revenue.head(10)

# Visualization of Top 10 Genres by Average Worldwide Gross
plt.figure(figsize=(14, 7))

# Use a seaborn color palette for aesthetics
sns.barplot(
    x=top_10_genres.index,
    y=top_10_genres.values,
    palette="Blues_d"
)

# Title and Labels
plt.title('Top 10 Genres by Average Worldwide Gross', fontsize=16, weight='bold')
plt.xlabel('Genre', fontsize=12)
plt.ylabel('Average Gross Revenue (in millions USD)', fontsize=12)

# Rotate x-axis labels and add gridlines
plt.xticks(rotation=30, fontsize=10)
plt.grid(axis='y', linestyle='--', alpha=0.7)


# Annotate bars with their values
for index, value in enumerate(top_10_genres.values):
    plt.text(index, value + (value * 0.01), f"${value:,.2f}", ha='center', fontsize=10)

plt.tight_layout()
plt.show()
```

```
genres
Animation      3.973884e+08
Adventure      3.796038e+08
Sci-Fi         3.715889e+08
Action         2.709311e+08
Fantasy        2.663832e+08
Family         2.000550e+08
Sport          1.801060e+08
Comedy         1.642282e+08
Western        1.524549e+08
Documentary    1.504499e+08
Musical        1.500816e+08
Thriller       1.437951e+08
History        1.155057e+08
Horror         1.101223e+08
Biography      1.093106e+08
Mystery        1.070858e+08
Crime          1.048086e+08
Music          1.036456e+08
Drama          1.023338e+08
War            9.268191e+07
Romance        9.153810e+07
News           6.283172e+07
Name: worldwide_gross, dtype: float64
```

```
C:\Users\USER\AppData\Local\Temp\ipykernel_10052\2325020745.py:18: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.



<Figure size 1000x600 with 0 Axes>

Budget vs. Worldwide Gross

How do budgets correlate with worldwide gross?

I'm interested in exploring how production budgets correlate with worldwide gross revenue. To investigate this, I'll use scatter plots to visualize the relationship and calculate correlation coefficients to quantify the strength of the association.


```
In [ ]: # Scatter plot of Production Budget vs Worldwide Gross
sns.scatterplot(
    data=merged_data,
    x='production_budget',
    y='worldwide_gross',
    alpha=0.6,
    edgecolor=None,
    color='blue'
)

# Add a regression line to visualize the trend
sns.regplot(
    data=merged_data,
    x='production_budget',
    y='worldwide_gross',
    scatter=False,
    color='red',
    line_kws={'linewidth': 2}
)

# Title and axis labels
plt.title('Relationship Between Production Budget and Worldwide Gross', fontsize=16)
plt.xlabel('Production Budget (in millions USD)', fontsize=12)
plt.ylabel('Worldwide Gross (in millions USD)', fontsize=12)

# Display the correlation coefficient on the plot
correlation = merged_data['production_budget'].corr(merged_data['worldwide_gross'])
plt.text(
    0.05, 0.95,
    f'Correlation: {correlation:.2f}',
    ha='left', va='top',
    transform=plt.gca().transAxes,
    fontsize=12, color='darkblue',
    bbox=dict(facecolor='white', alpha=0.7, edgecolor='blue')
)

# Show the plot
plt.tight_layout()
plt.show()
```



Profitable Release Windows

Q3: What are the most profitable release windows for movies?

I want to find out which release windows are the most profitable for movies. By analyzing the average revenue of films released in different months

```
In [48]: # Group the data by 'release_month' and calculate the average 'profit' for each mon
# Then, sort the resulting values in descending order to find the months with the h
release_profit = merged_data.groupby('release_month_name')['profit'].mean().sort_va
display(release_profit)

# Visualization of Average profit by release month
# Calculate average profit by release month
release_profit = merged_data.groupby('release_month_name')['profit'].mean().sort_va

# Set a consistent style
sns.set_style("whitegrid")

# Create the bar plot
plt.figure(figsize=(14, 8))
sns.barplot(
    x=release_profit.index,
    y=release_profit.values,
    palette="coolwarm"
)
```

```

# Add title and labels
plt.title('Average Profit by Release Month', fontsize=16, weight='bold')
plt.xlabel('Release Month', fontsize=12)
plt.ylabel('Average Profit (in millions USD)', fontsize=12)

# Rotate x-axis labels and add gridlines
plt.xticks(rotation=30, fontsize=10)
plt.grid(axis='y', linestyle='--', alpha=0.7)

# Annotate bars with their values
for index, value in enumerate(release_profit.values):
    plt.text(index, value + (value * 0.01), f"${value:,.2f}", ha='center', fontsize=10)

# Ensure proper layout
plt.tight_layout()
plt.show()

```

```


release_month_name
May                2.042898e+08
June               1.923356e+08
July               1.614964e+08
November           1.489081e+08
April              1.304827e+08
February           1.255192e+08
December           1.155203e+08
September          1.090012e+08
March              9.157966e+07
August             7.543899e+07
January            6.866001e+07
October            6.599796e+07
Name: profit, dtype: float64

```

C:\Users\USER\AppData\Local\Temp\ipykernel_10052\2865850718.py:16: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(
```



Studio Profit Margins

Q4: Which studios have the highest profit margins?

I'm curious about which studios achieve the highest profit margins.

To compare profitability—calculated as `worldwidegross` minus `budget`—across production studios

```
In [48]: # Group the data by 'studio' and calculate the average 'profit_margin' for each stu
# The 'profit_margin' is the ratio of profit to the production budget, which measur
studio_profit = merged_data.groupby('studio')['profit_margin'].mean()

# Sort the resulting series in descending order to find the studios with the highes
# This step ensures that the studios with the highest ROI are shown at the top
studio_profit = studio_profit.sort_values(ascending=False)

# Display the resulting series
# Display the top 10 studios with the highest profit margin
top_10_studios = studio_profit.head(10)

# Set a consistent style for the plot
sns.set_style("whitegrid")

# Visualization of top 10 studios by profit margin
plt.figure(figsize=(14, 8))

# Use seaborn's barplot for a cleaner look
sns.barplot(
    x=top_10_studios.index,
```

```

    y=top_10_studios.values,
    palette="viridis"
)

# Add title and axis labels
plt.title('Top 10 Studios by Profit Margin (ROI)', fontsize=16, weight='bold')
plt.xlabel('Studio', fontsize=12)
plt.ylabel('Average Profit Margin', fontsize=12)

# Rotate x-axis labels for readability and add gridlines
plt.xticks(rotation=30, fontsize=10)
plt.grid(axis='y', linestyle='--', alpha=0.7)


# Annotate bars with their values
for index, value in enumerate(top_10_studios.values):
    plt.text(index, value + 0.01, f"{value:.2f}", ha='center', fontsize=10)

# Adjust layout to prevent clipping
plt.tight_layout()
plt.show()

```

C:\Users\USER\AppData\Local\Temp\ipykernel_11128\395597088.py:20: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.



Language Influence on Global Performance

Q5: Does the original language influence global performance?

```
In [49]: # Group the data by original language and calculate average worldwide gross and count
language_stats = merged_data.groupby('original_language').agg(
    average_worldwide_gross=('worldwide_gross', 'mean'), # Calculate the average w
    number_of_movies=('worldwide_gross', 'count') # Count the number of mov
)

# Filter the data to include only languages with more than 2 movies
language_stats_filtered = language_stats[language_stats['number_of_movies'] > 2]

# Sort the filtered data by average worldwide gross
language_stats_filtered = language_stats_filtered.sort_values(by='average_worldwide_gross')

# Display the filtered and sorted dataframe
display(language_stats_filtered)

# Prepare the data for visualization
language_revenue = language_stats_filtered['average_worldwide_gross']

# Set a consistent style for the plot
sns.set_style("whitegrid")

# Create a bar plot for average worldwide gross by original language
plt.figure(figsize=(14, 8))
sns.barplot(
    x=language_revenue.index,
    y=language_revenue.values,
    palette="coolwarm"
)

# Add titles and labels
plt.title('Worldwide Gross by Original Language (Filtered by >2 Movies)', fontsize=14)
plt.xlabel('Original Language', fontsize=12)
plt.ylabel('Average Worldwide Gross (in millions USD)', fontsize=12)

# Rotate x-axis labels for better readability and add gridlines
plt.xticks(rotation=45, fontsize=10)
plt.grid(axis='y', linestyle='--', alpha=0.7)

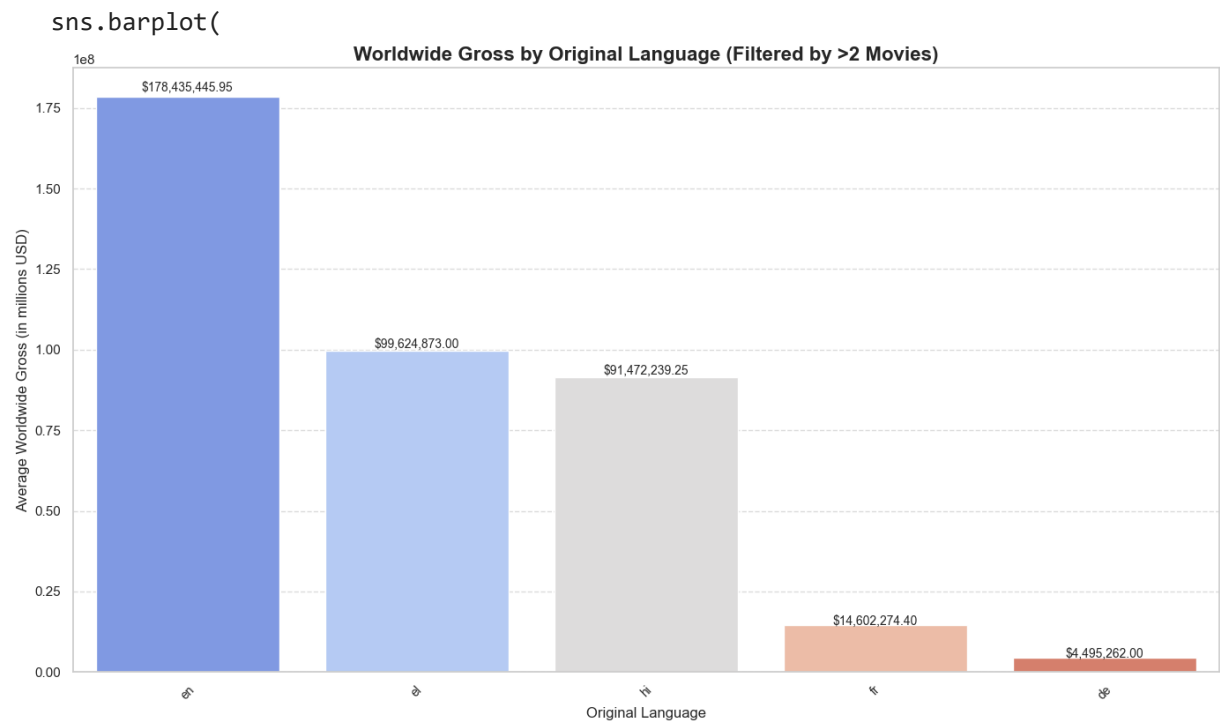
# Annotate bars with their values
for index, value in enumerate(language_revenue.values):
    plt.text(index, value + (value * 0.01), f"${value:,.2f}", ha='center', fontsize=10)

# Ensure layout doesn't get cut off
plt.tight_layout()
plt.show()
```


	average_worldwide_gross	number_of_movies
original_language		
en	1.784354e+08	1214
el	9.962487e+07	3
hi	9.147224e+07	4
fr	1.460227e+07	5
de	4.495262e+06	4

C:\Users\USER\AppData\Local\Temp\ipykernel_11128\2929811603.py:24: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.



Ratings and Revenue

Q6: How do ratings (average and votes) correlate with revenue?

```
In [50]: # Define colors for consistency
scatter_color = 'dodgerblue' # Blue for scatter points
regression_color = 'red'     # Red for regression line

# Scatter plot: Average Rating vs Worldwide Gross
plt.figure(figsize=(10, 6))
```

```
# Scatter plot for Average Rating vs Worldwide Gross with uniform color
sns.scatterplot(data=merged_data, x='averagerating', y='worldwide_gross', alpha=0.7)

# Add regression line with the same color scheme
sns.regplot(data=merged_data, x='averagerating', y='worldwide_gross', scatter=False)

# Title and Labels
plt.title('Average Rating vs. Worldwide Gross', fontsize=16, weight='bold')
plt.xlabel('Average Rating', fontsize=12)
plt.ylabel('Worldwide Gross (in millions USD)', fontsize=12)

# Add the correlation coefficient
correlation_rating_gross = merged_data['averagerating'].corr(merged_data['worldwide_gross'])
plt.text(0.05, 0.95, f'Correlation: {correlation_rating_gross:.2f}', ha='left', va='top')

# Show the plot
plt.tight_layout()
plt.show()

# Scatter plot: Number of Votes vs Worldwide Gross
plt.figure(figsize=(10, 6))

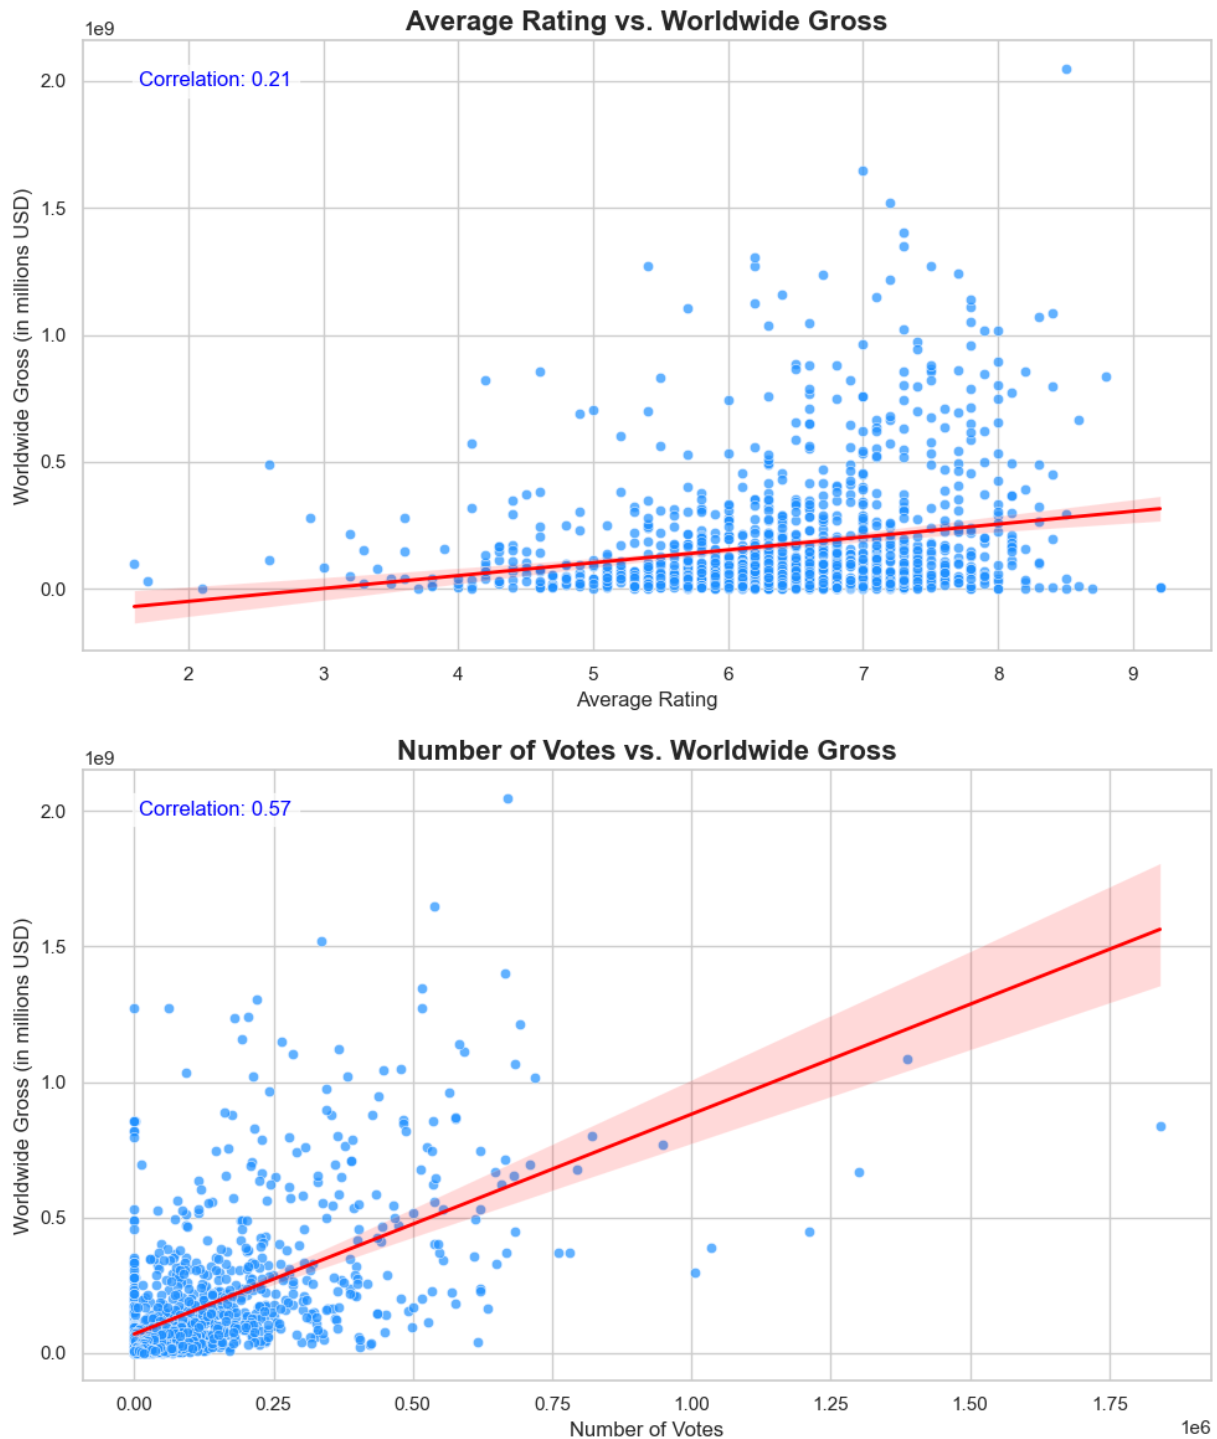
# Scatter plot for Number of Votes vs Worldwide Gross with uniform color
sns.scatterplot(data=merged_data, x='numvotes', y='worldwide_gross', alpha=0.7, color='blue')

# Add regression line with the same color scheme
sns.regplot(data=merged_data, x='numvotes', y='worldwide_gross', scatter=False, color='blue')

# Title and Labels
plt.title('Number of Votes vs. Worldwide Gross', fontsize=16, weight='bold')
plt.xlabel('Number of Votes', fontsize=12)
plt.ylabel('Worldwide Gross (in millions USD)', fontsize=12)

# Add the correlation coefficient
correlation_votes_gross = merged_data['numvotes'].corr(merged_data['worldwide_gross'])
plt.text(0.05, 0.95, f'Correlation: {correlation_votes_gross:.2f}', ha='left', va='top')

# Show the plot
plt.tight_layout()
plt.show()
```



4. Conclusion

Key Insights:

1. **Animation, Adventure, and Sci-Fi** genres generate the highest revenue, not Action and Adventure alone.
2. Higher production budgets generally lead to higher worldwide gross, but profitability requires strategic spending.
3. May and June (summer release windows) are the most profitable months for movie

launches.

4. Smaller studios like **UTV** can achieve higher profit margins through operational efficiency, countering the dominance of large-budget studios.
5. English-language movies dominate global revenue, but well-crafted non-English films can succeed in niche markets.
6. Audience engagement (measured by votes) positively impacts revenue, but high ratings alone don't guarantee success.

Recommendations:

- Focus on producing **Animation, Adventure**, and **Sci-Fi** movies for maximum revenue.
- Target summer and holiday release windows, particularly in May and June.
- Learn from efficient studios like **UTV** to optimize profit margins.
- Invest in audience engagement strategies to increase buzz and votes for films.